Black Lives Matter Sentiment Analysis

Francesco Di Salvo

Politecnico di Torino

s282417

s282414@studenti.polito.it

Abstract—In this report we will discuss a possible approach for a Twitter sentiment analysis. In particular this analysis regards the Black Lives Matter movement. This solution lays the foundation about the Natural Language Processing by adopting some of the most common techniques about text analysis, such as tokenization, stemming, stopwords elimination and n-grams implementation.

I. PROBLEM OVERVIEW

This sentiment analysis is based on a collection of 80.000 tweets collected from 2020-05-31 to 2020-06-08 about the Black Lives Matter movement. This movement is a decentralized political and social movement advocating for non-violent civil disobedience in protest against incidents of police brutality and all racially motivated violence against black people [1]. It was born in 2013 but it (sadly) becomes well known all over the word after the murder of George Floyd on 2020-05-25.

These tweets are already classified as:

- *1*, if they shows a positive sentiment towards the movement
- *0*, if they shows a negative sentiment towards the movement

Each tweet is described by 28 features and it is possible to see them on the Tweet Object Documentation [2]. The class label was added for the special purpose of the analysis. For this specific approach we focused just on the label and on the relative text, contained on the *full_text* feature.

This problem is well balanced, in fact we have 80.000 tweets, 40.069 classified as *0* and the remaining 39.931 as *1*. There was a strong activity during this week, in fact we have found in this sample a number of retweets equal to 981.031.288. As you can see in Figure 1, the daily activity increased until 2020-06-03 where it reached its peak and it strongly decreased until the last day to our disposal.

We have also analyzed the most frequent words inside both categories and you can see it in Figure 2. In particular they are both "coherent" with what we could expect from them. In fact for the "positive" tweets we can see words as "petitions", "everyone", "together", while for post against the movement there are words as "white", "division", "behind" and so on.

II. PROPOSED APPROACH

A. Data preprocessing

Natural Language Processing problems have a canonical preprocessing structure, in fact in order to analyze a collection of document we should consider several consecutive steps. In
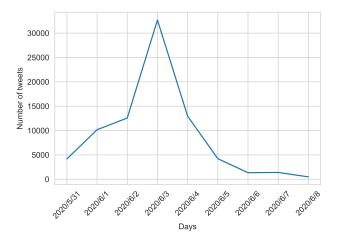

Fig. 1. Tweets per day

Fig. 2. Most frequent words on tweets supporting the cause

particular when we deal with huge documents, we can actually *split* them in shorter pieces in order to better analyze their content. Since this is not the case, we skipped to the second step, that is the *tokenization*: here we split each content into tokens (a.k.a words). These tokens will represent the "core" of our analysis, because some of these tokens will be strictly related to our class.

In particular we used the *Natural Language Tool Kit tokenizer*, that allows to easily convert a sentence like

Fig. 3. Most frequent words on tweets against the cause

"we are equal" into the following array of features ["we","are","equal"].

The aim of any sentiment analysis is to extract the most relevant features for each class. In this case, the most relevant ones, are the most frequent ones. In order to count the frequency of each token, we need to represents them as a *bag of words*, in particular we define a huge sparse matrix that stores counts of all the words in our collection of tweets. To do so, we used the fairly well known *TfIdfVectorizer* defined in scikit-learn. TF states for term frequencies and IDF states for inverse document frequency. It means that it penalize the tokens that are frequent in all the documents (our tweets), and it is defined as :

$$TF - IDF(t) = freq(t,d) \times \log \frac{m}{freq(t,D)} \quad (1)$$

where *t* is the token, *d* is the current document, and *D* is the collection of *m* documents. So, that's the basic idea of the sentiment analysis, but there are several other techniques that can help increase the performance. In particular we adopted the *stemming* and the *n-gram* implementation. The stemming comes from linguistic and it extracts the root form for each word, for example "consult, consulting, consultant" will all be treated as "consult", in order to unify analogous words. Then, the n-grams are all combinations of adjacent words or letters that can be extracted from the initial text. In particular we considered unigrams, bigrams and threegrams, by defining the range (1,3). To be more precise, we actually tried to remove also the *stopwords* that are the most common words present in the English grammar (articles, noun, pronoun and so on) but we noticed a decrease in performance. It is due to the fact that probably the stopwords that we used contained some relevant words for our analysis.

B. Model selection

The proposed approach has been tested on two classifiers:

- *logistic regression*: a linear model used to estimate the probability that an istance belongs to a particular class

TABLE I
HYPERPARAMETERS USED FOR THE GRIDSEARCHCV

Model	Hyperparameters	Values
LogisticRegressor	*C*	{1e-03, 1e-02, ... , 1e+03}
	penalty	{'l1','l2'}
MultinomialNB	*alpha*	{0.5,0.7,...,1.5}
	fit_prior	{True,False}

by computing a weighted sum of the input features and returns the logistic of its result,

- *multinomial naive bayes classifier*, a classifier that lays its foundations on the bayes theorem and on the assumptions of the mutually indipendent from the features. This classifier is fairly used in natural language processing because of its semplicity, in fact it is also much faster than the logistic regression.

A good compromise between accuracy and speed should be taken into account for business oriented application, because sometimes an accurate model could be so expensive for a real word application. Then, for both classifiers, the best hyperparameters configurations was defined through a grid search, explained in detail in the following section.

C. Hyperparameters tuning

There have been defined two sets of different yperparameters, defined in Table 1. The grid search trains the model with all the possible configurations and measure its performance by using the cross validation method. This method randomly split the entire dataset into k folds, train the model on k-1 folds and test it with the remaining one. This process is repeated until each fold has been tested. At the end of the process, it is possible to select the best configuration (the one that performs the best score) and use it for testing the evaluation test.

III. RESULTS

The reference metric was the accuracy and the best configuration for the LogisticRegression was {*'C':1000.0, 'penalty':'l2'*} obtained a score of 0.93. It outperformed the Multinomial Naive Bayes classifer that scored 0.906 with {*'alpha':0.5, 'fit_prior':True*}. Then, they both outperformed their equivalent naive solutions, where we just tokenized the corpus, avoidind sophisticated techniques as we explained before. It scored 0.926 with the LogisticRegression and 0.894 with the MultinomialNB, that are not so far from the our best one.

IV. DISCUSSION

The Natural Language Processing is a wide and interesting field, with several application on our lives, in fact we will make some others analysis in order to learn new things and to discover new interesting insights. Anyway, the current approach obtained an overall satisfactory result, and it probably also depends on the goodness of the provided dataset. Possible improvement could regard the preprocessing phase, by considering some other raffinate techniques. Then it should

be useful to focus on new models and their response to the different "combinations" of preprocessing.

REFERENCES

[1] Wikipedia "Black Lives Matter" Phil. Trans. Roy. Soc. London, vol. A247, pp. 529–551, April 1955.

[2] Giulio Angiani, Laura Ferrari, Tomaso Fontanini, Paolo Fornacciari, EleonoraIotti, Federico Magliani, and Stefano Manicardi "A Comparison between PreprocessingTechniques for Sentiment Analysis in Twitter"